

Assignment39

John Bute

2024-11-19

R Markdown

5. Construct and evaluate ANN models for the Wines dataset

```
library(nnet)
Wine <- read.csv("C:/Users/johnb/Desktop/Machine Learning/data/wine.csv",
stringsAsFactors = TRUE)

set.seed(1)
Wine$Class = as.factor(Wine$Class)

index <- sample(1:nrow(Wine), 0.7 * nrow(Wine))
train_data <- Wine[index, ]
test_data <- Wine[-index, ]

ann_model <- nnet(Class ~ ., data = train_data, size = 10, maxit = 100)

## # weights: 173
## initial value 215.368774
## iter 10 value 135.130137
## final value 135.130134
## converged

summary(ann_model)

## a 13-10-3 network with 173 weights
## options were - softmax modelling
## b->h1 i1->h1 i2->h1 i3->h1 i4->h1 i5->h1 i6->h1 i7->h1 i8->h1 i9->h1
## 0.54 0.00 0.53 -0.44 0.36 0.31 0.62 0.07 0.30 -0.16
## i10->h1 i11->h1 i12->h1 i13->h1
## -0.56 0.60 -0.30 0.13
## b->h2 i1->h2 i2->h2 i3->h2 i4->h2 i5->h2 i6->h2 i7->h2 i8->h2 i9->h2
## -0.55 0.48 -0.25 0.40 -0.33 -0.39 0.02 -0.32 -0.45 0.03
## i10->h2 i11->h2 i12->h2 i13->h2
## 0.09 -0.52 -0.34 0.31
## b->h3 i1->h3 i2->h3 i3->h3 i4->h3 i5->h3 i6->h3 i7->h3 i8->h3 i9->h3
## 0.65 -0.55 0.37 0.63 0.45 -0.16 0.21 0.64 0.64 -0.22
## i10->h3 i11->h3 i12->h3 i13->h3
## -0.32 -0.47 -0.25 1.99
## b->h4 i1->h4 i2->h4 i3->h4 i4->h4 i5->h4 i6->h4 i7->h4 i8->h4 i9->h4
## 0.59 0.02 -0.34 -0.63 -0.12 0.50 -0.21 -0.52 -0.18 0.18
## i10->h4 i11->h4 i12->h4 i13->h4
## -0.15 0.27 0.27 0.08
## b->h5 i1->h5 i2->h5 i3->h5 i4->h5 i5->h5 i6->h5 i7->h5 i8->h5 i9->h5
## -0.10 -0.07 -0.27 0.11 0.57 -0.50 -0.12 -0.40 -0.10 -0.51
## i10->h5 i11->h5 i12->h5 i13->h5
## -0.06 0.62 0.37 0.61
## b->h6 i1->h6 i2->h6 i3->h6 i4->h6 i5->h6 i6->h6 i7->h6 i8->h6 i9->h6
## -0.04 0.15 -0.02 -0.55 -0.35 0.00 -0.18 0.61 0.03 -0.26
## i10->h6 i11->h6 i12->h6 i13->h6
## -0.31 0.40 0.28 -0.47
## b->h7 i1->h7 i2->h7 i3->h7 i4->h7 i5->h7 i6->h7 i7->h7 i8->h7 i9->h7
## -0.61 0.36 0.17 -0.46 -0.61 -0.55 -0.17 -0.46 -0.28 -0.43
## i10->h7 i11->h7 i12->h7 i13->h7
```

```
##      -0.34    -0.45    -0.03     0.38
## b->h8 i1->h8 i2->h8 i3->h8 i4->h8 i5->h8 i6->h8 i7->h8 i8->h8 i9->h8
##      -0.66     0.04     0.53    -0.18    -0.63    -0.51    -0.25    -0.48    -0.51    -0.39
## i10->h8 i11->h8 i12->h8 i13->h8
##      -0.38    -0.52     0.67    -0.24
## b->h9 i1->h9 i2->h9 i3->h9 i4->h9 i5->h9 i6->h9 i7->h9 i8->h9 i9->h9
##      0.01     0.25    -0.56    -0.53    -0.63     0.60     0.24    -0.57    -0.01    -0.05
## i10->h9 i11->h9 i12->h9 i13->h9
##      -0.17     0.69    -0.45     0.44
## b->h10 i1->h10 i2->h10 i3->h10 i4->h10 i5->h10 i6->h10 i7->h10
##      -0.60    -0.14    -0.50    -0.43     0.48     0.31    -0.33    -0.01
## i8->h10 i9->h10 i10->h10 i11->h10 i12->h10 i13->h10
##      -0.58    -0.20     0.66     0.17     0.23    -0.26
## b->o1 h1->o1 h2->o1 h3->o1 h4->o1 h5->o1 h6->o1 h7->o1 h8->o1 h9->o1
##      -0.70     0.13    -0.07     2.42    -0.13    -0.17    -0.32    -0.20     0.68    -0.86
## h10->o1
##      -0.14
## b->o2 h1->o2 h2->o2 h3->o2 h4->o2 h5->o2 h6->o2 h7->o2 h8->o2 h9->o2
##      0.40    -0.63    -0.23     0.87    -0.52     0.08     0.66     0.65    -0.45     0.02
## h10->o2
##      -0.16
## b->o3 h1->o3 h2->o3 h3->o3 h4->o3 h5->o3 h6->o3 h7->o3 h8->o3 h9->o3
##      0.85     0.28     0.56    -3.19     0.42     0.92    -0.02    -0.01     0.40     0.49
## h10->o3
##      0.67
```

```
ann_preds <- predict(ann_model, newdata = test_data, type = "class")
conf_matrix_ann <- table(Actual = test_data$Class, Predicted = ann_preds)
print(conf_matrix_ann)
```

```
##      Predicted
## Actual  2
##        1 20
##        2 22
##        3 12
```

```
test_error <- 1 - sum(diag(conf_matrix_ann)) / sum(conf_matrix_ann)
cat("Test Error Rate:", test_error, "\n")
```

```
## Test Error Rate: 0.6296296
```

Well, it seems that it predicts just one class, and failing to identify any other class. It is highly biased and no generalization at all.

6. Re-run the ANN models incorporating 10 hidden layers with 30 nodes. How much more time does it take to run a “bigger” neural network on the Wine dataset? Let us compare both model’s run times:

```
library(nnet)
library(microbenchmark)

first_ann_benchmark <- microbenchmark(
  ann_model <- nnet(Class ~ ., data = train_data, size = 10, maxit = 100),
  times = 10L
)

## # weights: 173
## initial value 144.310736
## iter 10 value 123.614907
## iter 20 value 100.890702
## iter 30 value 80.875859
## iter 40 value 74.174068
## iter 50 value 60.066808
## iter 60 value 43.865311
```

```
## iter 70 value 24.922794
## iter 80 value 19.438731
## iter 90 value 15.428540
## iter 100 value 13.573137
## final value 13.573137
## stopped after 100 iterations
## # weights: 173
## initial value 188.594100
## final value 135.130134
## converged
## # weights: 173
## initial value 173.203244
## final value 135.130134
## converged
## # weights: 173
## initial value 156.889627
## final value 135.130133
## converged
## # weights: 173
## initial value 135.589664
## final value 135.130134
## converged
## # weights: 173
## initial value 147.344042
## final value 135.130134
## converged
## # weights: 173
## initial value 150.832781
## final value 135.130134
## converged
## # weights: 173
## initial value 140.709882
## final value 135.130134
## converged
## # weights: 173
## initial value 209.059997
## final value 135.130134
## converged
## # weights: 173
## initial value 147.833100
## iter 10 value 121.007508
## iter 20 value 96.293464
## iter 30 value 96.276626
## final value 96.276595
## converged
```

```
print(first_ann_benchmark)
```

```
## Unit: milliseconds
```

```
##
## ann_model <- nnet(Class ~ ., data = train_data, size = 10, maxit = 100)
## min lq mean median uq max neval
## 2.905001 3.004401 5.481871 3.180951 3.604901 21.9333 10
```

So, it takes about 13 milliseconds to run the smaller model. What about the bigger one

```
library(nnet)
```

```
library(microbenchmark)
```

```
big_ann_benchmark <- microbenchmark(
  ann_model <- nnet(Class ~ ., data = train_data, size = 30, maxit = 500),
  times = 10L
)

## # weights: 513
## initial value 236.471417
## final value 135.130134
```

```
## converged
## # weights:  513
## initial  value 239.262237
## final   value 135.130134
## converged
## # weights:  513
## initial  value 219.097102
## iter   10 value 131.872426
## iter   20 value 64.314334
## iter   30 value 60.431177
## iter   40 value 60.007864
## iter   50 value 57.829738
## iter   60 value 56.838110
## iter   70 value 56.757665
## iter   80 value 55.637476
## iter   90 value 55.356439
## iter  100 value 55.201881
## iter  110 value 54.092181
## iter  120 value 52.326306
## iter  130 value 52.018193
## iter  140 value 51.972187
## iter  150 value 51.767332
## iter  160 value 51.348508
## iter  170 value 51.185806
## iter  180 value 51.004440
## iter  190 value 47.748294
## iter  200 value 26.087473
## iter  210 value 9.441419
## iter  220 value 7.083028
## iter  230 value 4.409302
## iter  240 value 0.769179
## iter  250 value 0.009451
## iter  260 value 0.000605
## final   value 0.000085
## converged
## # weights:  513
## initial  value 390.963771
## iter   10 value 135.130134
## iter   10 value 135.130134
## iter   10 value 135.130134
## final   value 135.130134
## converged
## # weights:  513
## initial  value 141.371714
## iter   10 value 135.130140
## iter   20 value 103.961145
## iter   30 value 94.252781
## iter   40 value 94.222307
## final   value 94.222209
## converged
```

```

## # weights:  513
## initial  value 187.739133
## final  value 135.130134
## converged
## # weights:  513
## initial  value 188.967699
## iter   10 value 135.130134
## iter   10 value 135.130134
## iter   10 value 135.130134
## final  value 135.130134
## converged
## # weights:  513
## initial  value 322.617980
## iter   10 value 121.054351
## iter   20 value 68.800790
## iter   30 value 63.606752
## iter   40 value 61.971744
## iter   50 value 60.894231
## iter   60 value 59.264826
## iter   70 value 58.302398
## iter   80 value 57.930437
## iter   90 value 57.282789
## iter  100 value 56.092305
## iter  110 value 41.634976
## iter  120 value 16.695990
## iter  130 value 10.029500
## iter  140 value 5.437891
## iter  150 value 1.603264
## iter  160 value 0.109639
## iter  170 value 0.002275
## iter  180 value 0.000503
## iter  190 value 0.000239
## iter  200 value 0.000114
## final  value 0.000097
## converged
## # weights:  513
## initial  value 262.557645
## final  value 135.130134
## converged
## # weights:  513
## initial  value 159.330308
## final  value 135.130133
## converged

print(first_ann_benchmark)

## Unit: milliseconds
##
##   ann_model <- nnet(Class ~ ., data = train_data, size = 10, maxit = 100)

```

```
##      min      lq      mean    median      uq      max neval
## 2.905001 3.004401 5.481871 3.180951 3.604901 21.9333    10

print(big_ann_benchmark)

## Unit: milliseconds
##
##      ann_model <- nnet(Class ~ ., data = train_data, size = 30, maxit = 500)
##      min      lq      mean    median      uq      max neval
## 7.575001 8.698501 43.73688 9.478151 33.9459 192.1392    10
```

As we see, the mean time for the bigger model was 101 milliseconds, while for the smaller model, it was 14 milliseconds.

Let us see if the bigger model performs better

```
set.seed(123)
big_ann_model <- nnet(Class ~ ., data = train_data, size = 30, maxit = 500)

## # weights: 513
## initial value 148.254270
## iter 10 value 133.479661
## iter 20 value 94.329901
## iter 30 value 81.278339
## iter 40 value 26.537377
## iter 50 value 25.660925
## iter 60 value 25.540541
## iter 70 value 22.992134
## iter 80 value 22.496830
## iter 90 value 20.180725
## iter 100 value 20.105207
## iter 110 value 20.090700
## iter 120 value 20.016665
## iter 130 value 20.013450
## iter 140 value 19.982257
## iter 150 value 19.839588
## iter 160 value 19.832572
## iter 170 value 19.824955
## iter 180 value 19.823714
## iter 190 value 19.823241
## iter 200 value 19.821929
## iter 210 value 19.820938
## iter 220 value 19.820600
## iter 230 value 15.564400
## iter 240 value 15.316459
## iter 250 value 15.293021
## iter 260 value 15.267555
## iter 270 value 15.228925
## iter 280 value 15.219195
## iter 290 value 15.218384
## iter 300 value 15.218296
```

```

## iter 310 value 15.217997
## iter 320 value 15.216966
## iter 330 value 15.210446
## iter 340 value 15.169683
## iter 350 value 15.125378
## iter 360 value 14.836218
## iter 370 value 14.610837
## iter 380 value 10.957461
## iter 390 value 10.654106
## iter 400 value 10.640421
## iter 410 value 10.520044
## iter 420 value 10.465496
## iter 430 value 10.366535
## iter 440 value 10.217303
## iter 450 value 10.047114
## iter 460 value 10.031337
## iter 470 value 10.031141
## iter 480 value 10.031033
## iter 490 value 10.030555
## iter 500 value 9.879454
## final value 9.879454
## stopped after 500 iterations

big_ann_preds <- predict(big_ann_model, newdata = test_data, type = "class")
conf_matrix_big_ann <- table(Actual = test_data$Class, Predicted =
big_ann_preds)
print(conf_matrix_big_ann)

##          Predicted
## Actual   1   2   3
##          1 19   1   0
##          2   2 18   2
##          3   1   0 11

test_error <- 1 - sum(diag(conf_matrix_big_ann)) / sum(conf_matrix_big_ann)
cat("Test Error Rate:", test_error, "\n")

## Test Error Rate: 0.1111111

```

We achieve an error rate of 11%, which is certainly way better than what we previously got.